

# How to develop with Kubernetes

## Developer workstation's tools

Damien Ribeiro



February 4, 2020

# Outline

- 1 First worker
- 2 Working in Kubernetes
- 3 Managing different configurations
- 4 Batch and scheduling
- 5 Web service

# Pre-requirements

- a little of bash-fu
- a computer with docker installed ([official documentation](#))

# Terminals

During the workshop, we will need to use several shell at the same time.

We need a tool for that, like for example:

- tmux (in command line)
- screen
- Terminator
- ... or just open several terminals

# Terminals

In this workshop we will use up to four terminals at once.

To keep things clear, these terminals will be named:

- Command terminal
- Monitor terminal
- Logs terminal
- Second command terminal

If the command to run needs to be in a specific folder, it will be indicated as well.

## Folder tree

For the following, all path will be assume to be under a *training* folder.

For example, assuming the path of *training* is */training*,  
the mention *folder/script.sh*  
will be equivalent to */training/folder/script.sh*.

# Using workstation tools

We need to install the tools used during the workshops:

<b>Docker</b>	<a href="#">Official release page</a>
<b>Kubectl</b>	<a href="#">Official installation documentation</a>
<b>Minikube</b>	<a href="#">Official installation documentation</a>
<b>Kustomize</b>	<a href="#">Official installation documentation</a>
<b>Scaffold</b>	<a href="#">Official installation documentation</a>
<b>Stern</b>	<a href="#">Official installation documentation</a>

## Warning about Kustomize and Skaffold

As it is now, Kustomize and Skaffold are still under construction. These tools give a great help but the version evolve quickly and the files version can change. To avoid compatibility issues, there is a need to define the update policy for those tool's version.

*As these problematic depends on the organisation of the team or company, it will not be talked here.*



## Using minikube

Minikube consist of a single node kubernetes cluster. Here a few commands useful for the workshop:

Get the cluster IP. This IP will be used as the one of each kubernetes servers (master or node):

```
minikube ip
```

For the registry, we will use the docker local registry of minikube.

## Using minikube

For the following, make sure that minikube is stopped:

Command terminal

```
minikube stop
```

Check minikube status:

Command terminal

```
minikube status
```

# Objectives

- Our first task is to create a wonderful worker that will keep telling us that it is alive.
- To keep things simple, this worker will be a shell script.
- Then we will run this worker in a container on our local docker installation.

# Creating a simple worker

First we place ourselves in the folder "worker". And we create the worker script:

worker/worker.sh

```
#!/bin/sh
while true;do
    echo "I'm alive and it is $(date)!"
    sleep 2
done
```

## Testing our worker

### Command terminal in folder worker

```
chmod u+x worker.sh  
./worker.sh
```

We can kill our worker with hitting **Ctrl+C**.

# Embedding our worker in a container

We need to write a Dockerfile:

worker/Dockerfile

```
FROM alpine
```

```
COPY worker.sh .
```

```
CMD /worker.sh
```

# Creating our image

We create the image of our worker:

Command terminal in folder worker

```
docker build -t tinkou/worker:v1 .
```

This image can be seen locally in docker:

Command terminal

```
docker images
```

# Running our container

Now that we have our worker image, it is time to run it in docker:

## Command terminal

```
docker run tinkou/worker:v1
```

We can see our worker output.

We can kill our container with **Ctrl+C**.



## Playing a little with our container

We can start our container in background and give it a name:

### Command terminal

```
docker run --name worker -d tinkou/worker:v1
```

We can still take a look to the logs:

### Command terminal

```
docker logs -f worker
```

We can quit the follow with **Ctrl+C**.

## Playing a little with our container

We can start, stop and see the logs of our container at will:

### Command terminal

```
docker ps  
docker stop worker  
docker ps  
docker start worker  
docker ps
```

This is always the same container running. We can see all existing containers to check it:

### Command terminal

```
docker ps -a
```

# Cleaning up

We stop and remove the containers and then remove the images  
(*beware it will clean everything*):

## Command terminal

```
docker stop $(docker ps -q)
docker rm $(docker ps -aq)
docker rmi $(docker images -q)
```

We can check that nothing remain with:

## Command terminal

```
docker ps
docker ps -a
docker images
```

## Going further...

To learn more about docker, you can check [docker official documentation](#) or use the docker built-in help:

### Command terminal

```
docker help
docker help build
...
```

Questions?

# Objectives

Now that we have a worker, we want to run it in a kubernetes cluster.

The first step will be to run a container into kubernetes to make sure that we can do it.

Then, we will run our worker in kubernetes.

# Using minikube

For the following, make sure that minikube is started:

Command terminal

```
minikube start
```

Check minikube status:

Command terminal

```
minikube status
```

# Accessing a Kubernetes cluster

To access a kubernetes cluster, we use kubectl client configured with the file `$home/.kube/config`

More information [here](#).

As we are using Minikube, starting it already configures kubectl.

As indicated in [this documentation](#), you can enable the completion for kubectl.



# Running our first container in the cluster

To begin, we are just going to deploy a container sending a ping:

## Command terminal


```
kubectl run pingpong --image=alpine ping 1.1.1.1
```

```
kubectl get deployments -o wide
```

```
kubectl get replicaset -o wide
```

```
kubectl get pods -o wide
```

# Kubernetes application architecture



# Preparing the environment

Starting from now, we are going to use 3 shells in 3 different windows/panes.

# Looking at logs in a kubernetes cluster

We want to display the logs of this container:

Logs terminal

```
kubect1 logs -f pingpong-XXXX-XXXX
```

Something happens and we need to recreate the pod:

Monitor terminal

```
kubect1 get pods -w
```

Commands terminal

```
kubect1 delete pod pingpong-XXXX-XXXX
```

# Looking at logs in a kubernetes cluster

There are several problems:

- we do not see which logs comes from which pod
- in fact we do not see the new containers logs

How to solve this and be able to follow logs even if pods are moving?

# Using stern to look at logs

So, we are going to redo the operation but this time using stern:

Command terminal

```
stern pingpong
```

Monitor terminal

```
watch kubectl get all
```

Second command terminal

```
kubectl delete pod pingpong-XXXX-XXXX
```

# Cleaning up the namespace

Before returning to our worker, a little cleaning can be wise:

## Command terminal

```
kubectl delete deployment pingpong  
kubectl get all
```

And stop the commands in the other windows.

# Running our worker in kubernetes

Now that we know how to run something in kubernetes, we are going to do it with our worker:

Logs terminal

```
stern worker
```

Monitor terminal

```
kubectl get pods -w
```

Command terminal

```
docker images
```

```
kubectl run worker --image=tinkou/worker:v1
```



# Troubleshooting

Kubectrl get pods indicate that something went wrong.  
Check kubernetes object to find the problem root cause:

## Command terminal

```
kubectrl describe deployment worker  
kubectrl describe replicaset worker-XXXX  
kubectrl describe pod worker-XXXX-XXXX
```

The "Events" are the interesting part here.

# Using a registry

The events indicates that the image can not be found.

The docker daemon in minikube is not the same as the local one.

We need to:

- link the local docker CLI to the minikube docker daemon
- create the image in the minikube docker daemon locale registry

## Command terminal

```
eval $(minikube docker-env)
docker images
docker build -t tinkou/worker:v1 .
```

# Using a registry

After a moment, the pod and stern starts displaying logs. What happened?

# Using a registry

After a moment, the pod and stern starts displaying logs. What happened?

The different elements of kubernetes retry periodically. And as the image is now available, the pods can successfully start.

First worker

## Working in Kubernetes

Managing different configurations

Batch and scheduling

Web service

Objectives

Configuring the environment

Running a container in kubernetes

Running our worker in kubernetes

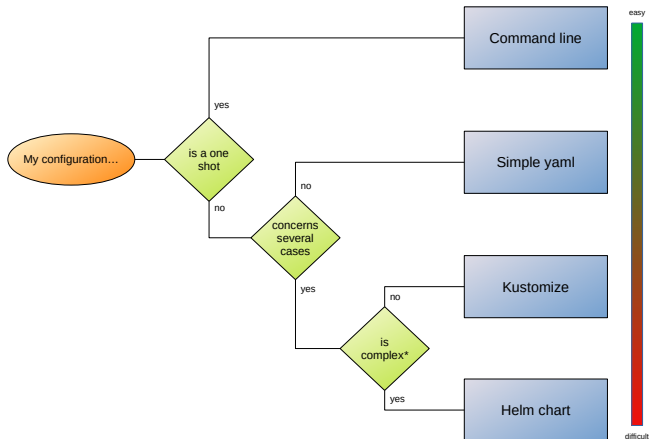
Questions?

# Objectives

Now that we have a worker in kubernetes, we want to be able to configure it.

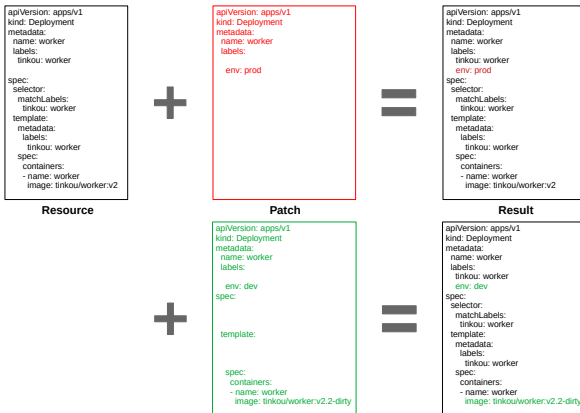
Also, we want to be able to differentiate between development and production configuration.

# Kustomize among other solutions



\* complex as in needing conditionals to valorize or depending on contextual values

# Kustomize: a layered base configuration





# Initializing the base

We create a folder tree to sort our files

Command terminal in folder worker

```
mkdir -p kube/base  
cd kube/base
```

Logs terminal still had "stern worker" running

Monitor terminal still had "kubectl get pods -w" running

# Initializing the base

We are starting with a source.yaml file describing our deployment

## worker/kube/base/source.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: worker
  labels:
    tinkou: worker
spec:
  selector:
    matchLabels:
      tinkou: worker
  template:
    metadata:
      labels:
        tinkou: worker
    spec:
      containers:
        - name: worker
          image: tinkou/worker:v1
```

# Initializing the base

Command terminal in folder worker/kube/base

```
kustomize create --autodetect  
kustomize build  
kubectl apply -k .
```

The command create will create a file kustomization.yaml in the current folder.

## Initializing the base

As we modify an immutable field, we need to delete the previous deployment before:

Command terminal in folder worker/kube/base

```
kubectl delete deployment worker  
kubectl apply -k .
```

## Adding a parameter to our worker

Modify the worker to use a parameter:

worker/worker.sh

```
#!/usr/bin/env bash
while true; do
    echo "I'm $NAME and it is $(date)"
    sleep 2
done
```

## Deploying the new version in our cluster

First we create a new version of the image:

Command terminal in folder worker

```
docker build -t tinkou/worker:v2 .
```

# Deploying the new version in our cluster

Update source.yaml the field `.spec.template.spec`:

worker/kube/base/source.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: worker
  labels:
    tinkou: worker
spec:
  selector:
    matchLabels:
      tinkou: worker
  template:
    metadata:
      labels:
        tinkou: worker
    spec:
      containers:
        - name: worker
          image: tinkou/worker:v2
          env:
            - name: NAME
              value: Bink
```

## Deploying the new version in our cluster

And finally apply the modification:

Command terminal in folder worker

```
kubectl apply -k kube/base
```



## Deploying the new version in our cluster

And finally apply the modification:

Command terminal in folder worker

```
kubectl apply -k kube/base
```

To clean what has been deployed:

Command terminal in folder worker

```
kubectl delete -k kube/base
```

# Deploying the new version in our cluster

There are too many operations.

Is there a way to simplify this?

# Skaffold

Skaffold is a developer oriented tool created to simplify the packaging and deployment on kubernetes.

The Skaffold documentation can be found [here](#).

# Initializing skaffold

Let's initialize skaffold:

Command terminal in folder worker

```
skaffold init
```

Follow the application command line interface.

This command create a file skaffold.yaml.

# Initializing scaffold

Scaffold detects the yaml and scaffold.yaml needs to be configured to use kustomize:

## worker/scaffold.yaml

```
apiVersion: scaffold/v1beta15
kind: Config
metadata:
  name: worker
build:
  artifacts:
    - image: tinkou/worker
deploy:
  kustomize:
    path: kube/base
```

The apiVersion can change with scaffold version.

# Using Skaffold

Let's use Skaffold to build our image:

Command terminal in folder worker

```
skaffold run
```

# Using Scaffold

Let's use Scaffold to build our image:

Command terminal in folder worker

```
scaffold run
```

A particularity of Scaffold is that to work with minikube, the current context of kubectl must be "minikube".

The current context can be with:

Command terminal

```
kubectl config current-context
```

# Using skaffold

We are going to test a little skaffold commands:

## Command terminal in folder worker

```
skaffold build  
skaffold run  
skaffold delete
```

**build** only build the image and push it in the registry (except in minikube)

**run** build, push and run the image in kubectl current cluster

**delete** delete run element from the kubectl current cluster



# Using skaffold

Lets try the dev function:

```
Command terminal in folder worker  
skaffold dev
```

It works like run, but displays logs and stays in foreground.

# Using scaffold

Open a command line in a new window 4.

Try to modify the file `worker.sh` in this new command line.

And look at how scaffold dynamically builds and deploys each time the file is saved.

# Isolating the configuration in a configmap

Add a file conf.env in the folder base:

```
worker/kube/base/conf.env
```

```
NAME=Trent
```

Second command terminal in the folder worker/kube/base

```
kustomize edit add configmap worker \  
--from-env-file conf.env
```

# Isolating the configuration in a configmap

## worker/kube/base/source.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: worker
  labels:
    tinkou: worker
spec:
  selector:
    matchLabels:
      tinkou: worker
  template:
    metadata:
      labels:
        tinkou: worker
    spec:
      containers:
        - name: worker
          image: tinkou/worker:v2
          envFrom:
            - configMapRef:
                name: worker
```

## Defining a local configuration

Create a folder `worker/kube/local`, and add in it the file `patch.yaml`:

`worker/kube/local/patch.yaml`

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: worker
data:
  NAME: Iris
```

## Defining a local configuration

Create a folder `worker/kube/local`, and add in it the file `patch.yaml`:

`worker/kube/local/patch.yaml`

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: worker
data:
  NAME: Iris
```

Nothing happens in scaffold dev.

## Defining a local configuration: kustomize

First create a new kustomize project with a base resource and the patch:

Second command terminal in folder worker/kube/local

```
kustomize create --resources ../base  
kustomize edit add patch patch.yaml
```

## Defining a local configuration: kustomize

First create a new kustomize project with a base resource and the patch:

Second command terminal in folder worker/kube/local

```
kustomize create --resources ../base  
kustomize edit add patch patch.yaml
```

Still nothing in scaffold dev.



## Defining a local configuration: skaffold

Indicate to skaffold to use this local values when using minikube:

### worker/skaffold.yaml

```
apiVersion: skaffold/v1beta15
kind: Config
metadata:
  name: worker
build:
  artifacts:
    - image: tinkou/worker
deploy:
  kustomize:
    path: kube/base
profiles:
- name: local
  activation:
    - kubeContext: minikube
  deploy:
    kustomize:
      path: /kube/local
```

## Using scaffold dev modifications detection

Now the modification is deployed.  
Lets modify the base configuration:

```
worker/kube/base/conf.env
```

```
NAME=Dor
```

## Using scaffold dev modifications detection

Now the modification is deployed.  
Lets modify the base configuration:

```
worker/kube/base/conf.env
```

```
NAME=Dor
```

Scaffold do not react.

Try to modify several files to see how scaffold detection works.

## Cleaning up

If you kill the command `skaffold dev` with **Ctrl+C**, skaffold will clean what it has deployed.

Questions?

# Objectives

Our worker is fine, but induces too much stress on the system. We now want to deploy it as a batch.

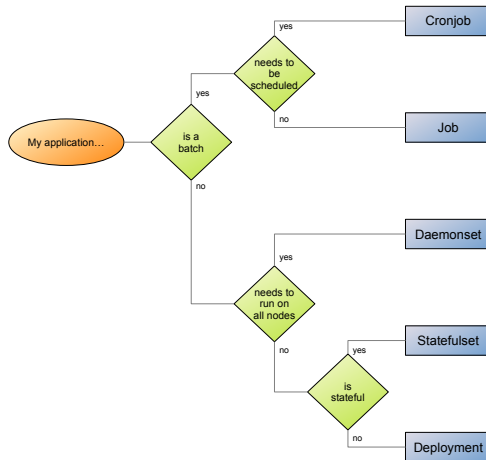
As such, we want to be able to schedule it.

We are now going to create a batch performing the same task as our worker.

## Objectives

Creating a batch in kubernetes  
Using a layer based configuration

# What kind to use to deploy our application



## Creating the batch

Beside the folder `worker`, create a folder `batch`.

Create the batch itself:

```
batch/batch.sh
```

```
echo "I'm Humphrey and it is $(date)"
```

Now, each execution is unitary.



## Exercise

Build and run in docker an image `tinkou/batch:v1` containing this batch.

# A solution

## batch/Dockerfile

```
FROM alpine

COPY batch.sh .
RUN chmod u+x batch.sh

CMD /batch.sh
```

## Command terminal in folder batch

```
docker build -t tinkou/batch:v1 .
docker run tinkou/batch:v1
```

# Creating our first CronJob

## batch/cronjob.yaml

```
apiVersion: batch/v1beta1
kind: CronJob
metadata:
  name: my-batch
spec:
  schedule: "* * * * *"
  jobTemplate:
    spec:
      template:
        metadata:
          labels:
            tinkou: batch
        spec:
          containers:
            - name: runner
              image: tinkou/batch:v1
          restartPolicy: Never
```

# Creating our first CronJob

## Logs terminal

```
stern -l tinkou=batch
```

## Monitor terminal

```
watch kubectl get all
```

## Command terminal

```
kubectl apply -f cronjob.yaml
```

Wait for a few batches to run...

```
kubectl logs -l tinkou=batch
```

# Creating our first CronJob

## Conclusions

- CronJob are easy to defined
- stern is less adapted than a standard kubectl logs for jobs

## CronJobs useful options

- A field `.spec.suspend` suspend the scheduling of a CronJob
- Jobs in error aren't removed
- The history limit of successful jobs can be set

## Exercise

Modify the configuration to use kustomize and scaffold with:

- a base configuration similar as the current one
- a local configuration for minikube with NAME=Dolph

# A solution

## Folder tree in batch folder

```
batch
|- batch.sh
|- Dockerfile
|- kube
  |- base
    |- conf.env
    |- cronjob.yaml
    |- kustomization.yaml
  |- local
    |- conf.yaml
    |- kustomization.yaml
```

batch/kube/base/conf.env

NAME=Humphrey

# A solution

## batch/kube/base/cronjob.yaml

```
apiVersion: batch/v1beta1
kind: CronJob
metadata:
  name: my-batch
spec:
  schedule: "* * * * *"
  jobTemplate:
    spec:
      template:
        metadata:
          labels:
            tinkou: batch
        spec:
          containers:
            - name: runner
              image: tinkou/batch:v1
              envFrom:
                - configMapRef:
                    name: batch
              restartPolicy: Never
```



## A solution

Command terminal in folder batch/kube/base

```
kustomize create --resources cronjob.yaml  
kustomize edit add configmap batch \  
                                --from-env-file conf.env
```

# A solution

batch/kube/local/conf.yaml

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: batch
data:
  NAME: Dolph
```

## A solution

Command terminal in folder batch/kube/local

```
kustomize create --resources ../base  
kustomize edit add patch conf.yaml
```

# A solution

## batch/skaffold.yaml

```
apiVersion: skaffold/v1beta12
kind: Config
build:
  artifacts:
    - image: tinkou/batch
deploy:
  kustomize:
    path: kube/base
profiles:
  - name: local
    activation:
      - kubeContext: minikube
    deploy:
      kustomize:
        path: kube/local
```

## A solution

Command terminal in folder batch

```
scaffold run
```

Take a quick look at the scheduling options and logs:

Command terminal

```
kubectl describe cronjob my-batch
```

Clean after working:

Command terminal in folder batch

```
scaffold delete
```

Questions?

# Objectives

Our project now needs a web service.

This web service needs to be monitored and scalable.

# Creating the web service itself

Create a new working folder web.

web/index.html

```
<!DOCTYPE HTML>
<html>
  <head>
    <meta charset="UTF-8">
    <title>Web service training</title>
  </head>
  <body>
    <h1>Hello World!</h1>
  </body>
</html>
```



# Creating the web service image

## web/Dockerfile

```
FROM nginx:alpine  
COPY index.html /usr/share/nginx/html
```

## Command terminal in folder web

```
docker build -t tinkou/web:v1 .
```

# Creating kubernetes deployment configuration

## web/deployment.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: web
  labels:
    tinkou: web
spec:
  selector:
    matchLabels:
      tinkou: web
  template:
    metadata:
      labels:
        tinkou: web
    spec:
      containers:
        - name: service
          image: tinkou/web:v1
```

# Creating kubernetes deployment configuration

## Logs terminal

```
stern -l tinkou=web
```

## Monitor terminal

```
watch kubectl get all
```

## Command terminal in folder web

```
kubectl apply -f deployment.yaml
```

# Creating kubernetes deployment configuration

A deployment, a replicaset and a pod have been created.

But how to access this web service?

## Exposing a web service

Pods have an IP...

Command terminal

```
kubectl get pod -o wide
```

... but this IP is internal to the cluster.

# Exposing a web service

Pods have an IP...

Command terminal

```
kubectl get pod -o wide
```

... but this IP is internal to the cluster.

—

In kubernetes, pods are exposed via services.

# Exposing a web service

## web/service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
  labels:
    tinkou: web
spec:
  type: NodePort
  selector:
    tinkou: web
  ports:
    - port: 80
      protocol: TCP
      nodePort: 32001
```

## Exposing a web service

### Command terminal in folder web

```
kubectl apply -f service.yaml  
curl http://$(minikube ip):32001
```

Now we can access the service.

It is also possible to display it in a web browser (by replacing `$(minikube ip)` by its value).



# How does it work?

- the minikube VM has its IP given by `minikube ip`
- as a NodePort, the service listen to the port 32001 of all nodes
- the service redirects requests in round robin to pods matching the selector
- pods listen to the same ports as their containers
- at last the container can answer to requests

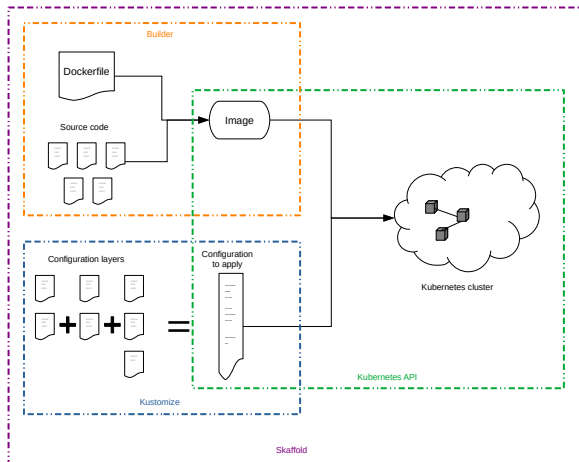
# Cleaning up

## Command terminal in folder web

```
kubectl delete -f deployment.yaml  
kubectl delete -f service.yaml
```

Questions?

# Conclusion



## Going further?

With these tools it is possible:


- to use an other image builder than docker (jib, kaniko, ...)
- to have stacks with multiple images
- to use any kind of kubernetes cluster instead of minikube
- to debug inside deployed containers

## A little more complex example

In [this github repository](#), you can find an example consisting of a complete stack.

This is an educational project based on [the work of jpetazzo](#).

## A little more complex example



## A little more complex example

This project is useful to demonstrate the usage of several languages and to manipulate kubernetes configuration.

Finally, in the branch *exercise*, there is the base configuration and it is possible to use it to try to create kustomize and skaffold configurations.



Questions?